



Estimating gene and transcript abundance using RNA-seq

Transcript Quantification: An Overview






10s-100s of millions of
short (35-300 character) "reads"



Given: (1) Collection of RNA-Seq fragments
(2) A set of known (or assembled) transcript sequences

Estimate: The relative abundance of each transcript



Sample

isoform A


isoform B

isoform C

% Gene 1

% Gene M


Abundance Estimates



10s-100s of millions of
short (35-300 character) "reads"

Given: (1) Collection of RNA-Seq fragments
(2) A set of **known** (or assembled) transcript sequences

Estimate: The relative abundance of each transcript



Sample

isoform A

isoform B

isoform C

% Gene 1

% Gene M

Abundance Estimates

Why not simply “count” reads


The RNA-seq reads are drawn from transcripts, and our spliced-aligners let us map them back to the transcripts on the genome from which they originate.

Problem: How do you handle reads that align equally-well to multiple isoforms / or multiple genes?

- Discarding multi-mapping reads leads to incorrect and biased quantification
- Even at the gene-level, the transcriptional output of a gene should depend on what isoforms it is expressing.

First, consider this non-Biological example

Imagine I have two colors of circle, **red** and **blue**. I want to estimate the **fraction of circles** that are **red** and **blue**. I'll *sample* from them by tossing down darts.




Here, a dot of a color means I hit a circle of that color.

What type of circle is more prevalent?

What is the fraction of red / blue circles?

First, consider this non-Biological example

Imagine I have two colors of circle, **red** and **blue**. I want to estimate the **fraction of circles** that are **red** and **blue**. I'll *sample* from them by tossing down darts.



You're missing a **crucial piece of information!**

The areas!


First, consider this non-Biological example

Imagine I have two colors of circle, **red** and **blue**. I want to estimate the **fraction of circles** that are **red** and **blue**. I'll *sample* from them by tossing down darts.

You're missing a **crucial piece of information!**


The areas!

There is an analog in RNA-seq, one needs to know the **length** of the target from which one is drawing to meaningfully assess abundance!




Resolving multi-mapping is fundamental to quantification

These errors can affect DGE calls



Resolving multi-mapping is fundamental to quantification

Can even affect abundance estimation in **absence** of alternative-splicing
(e.g. paralogous genes)




Main challenges of fast & accurate quantification

- finding locations of reads (alignment) is slower than necessary → simply aligning reads in a sample **can take hours**
- **alternative splicing** and **related sequences** creates ambiguity about where reads came from → **multi-mapping reads** *cannot* be ignored / discarded or assigned naïvely
- **sampling of reads is not uniform or idealized, exhibits multiple types of bias** → RNA-seq can exhibit ***extensive and sample-specific bias***
- uncertainty in ML estimate of abundances → There is both technical (shot noise) and ***inherent inferential uncertainty*** in abundance estimates

How can we perform inference from sequenced fragments?

Experimental Mixture




In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type

How can we perform inference from sequenced fragments?

Experimental Mixture




length() = 100

In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type

How can we perform inference from sequenced fragments?

Experimental Mixture




In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type

length() = 100 x 6 copies

How can we perform inference from sequenced fragments?

Experimental Mixture




In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type

$$\text{length}(\text{—————}) = 100 \times 6 \text{ copies} = 600 \text{ nt}$$

How can we perform inference from sequenced fragments?

Experimental Mixture



In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type


length() = 100 x 6 copies = 600 nt

length() = 66 x 19 copies = 1254 nt

length() = 33 x 6 copies = 198 nt

How can we perform inference from sequenced fragments?

Experimental Mixture




In an unbiased experiment, sampling fragments depends on:

- # of copies of each txp type
- length of each txp type

length() = 100	x 6 copies	= 600 nt	~ 30% blue
length() = 66	x 19 copies	= 1254 nt	~ 60% green
length() = 33	x 6 copies	= 198 nt	~ 10% red

How can we perform inference from sequenced fragments?

Experimental Mixture



In an unbiased experiment, sampling fragments depends on:


- # of copies of each txp type
- length of each txp type

length() = 100	x 6 copies	= 600 nt	~ 30% blue
length() = 66	x 19 copies	= 1254 nt	~ 60% green
length() = 33	x 6 copies	= 198 nt	~ 10% red

We call these values $\eta = [0.3, 0.6, 0.1]$ the nucleotide fractions,
they become the primary quantity of interest

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?

Think about the “ideal” RNA-seq experiment . . .




(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

How can we perform inference from sequenced fragments?


Think about the “ideal” RNA-seq experiment . . .



(1) Pick transcript $\mathbf{t} \propto$ total available nucleotides = count * length

(2) Pick a position $\mathbf{p}$ on $\mathbf{t}$ “uniformly at random”

Resolving a single multi-mapping read



Say we *knew* the η , and observed a *single* read that mapped ambiguously, as shown above.

What is the probability that it truly originated from **G** or **R**?

$$\Pr \{r \text{ from } G\} = \frac{\frac{\eta_G}{\text{length}(G)}}{\frac{\eta_G}{\text{length}(G)} + \frac{\eta_R}{\text{length}(R)}} = \frac{\frac{0.6}{66}}{\frac{0.6}{66} + \frac{0.1}{33}} = 0.75$$

$$\Pr \{r \text{ from } R\} = \frac{\frac{\eta_R}{\text{length}(R)}}{\frac{\eta_G}{\text{length}(G)} + \frac{\eta_R}{\text{length}(R)}} = \frac{\frac{0.1}{33}}{\frac{0.6}{66} + \frac{0.1}{33}} = 0.25$$

normalization factor

length() = 100 x 6 copies = 600 nt ~ 30% **blue**

length() = 66 x 19 copies = 1254 nt ~ 60% **green**

length() = 33 x 6 copies = 198 nt ~ 10% **red**

Units for Relative Abundance

TPM (Transcripts Per Million)

$$\text{TPM}_i = \rho_i \times 10^6 \text{ where } 0 \leq \rho_i \leq 1 \text{ and } \sum_i \rho_i = 1$$

$$\rho_i = \frac{\frac{X_i}{\ell_i}}{\sum_j \frac{X_j}{\ell_j}}$$

← Reads coming from transcript i

Units for Relative Abundance

TPM (Transcripts Per Million)

$$\text{TPM}_i = \rho_i \times 10^6 \text{ where } 0 \leq \rho_i \leq 1 \text{ and } \sum_i \rho_i = 1$$

$$\rho_i = \frac{\frac{X_i}{\ell_i}}{\sum_j \frac{X_j}{\ell_j}}$$

Reads coming from transcript i

Length of transcript i

Units for Relative Abundance

TPM (Transcripts Per Million)

$$\text{TPM}_i = \rho_i \times 10^6 \text{ where } 0 \leq \rho_i \leq 1 \text{ and } \sum_i \rho_i = 1$$

abundance of i
as fraction of all
measured transcripts

$$\rho_i = \frac{\frac{X_i}{\ell_i}}{\sum_j \frac{X_j}{\ell_j}}$$

Reads coming from
transcript i

Length of transcript i

A probabilistic view of RNA-Seq quantification

nucleotide fractions known transcriptome assumes independence of fragments

$$\Pr\{\mathcal{F} \mid \boldsymbol{\eta}, \mathcal{T}\} = \prod_{j=1}^N \Pr\{f_j \mid \boldsymbol{\eta}, \mathcal{T}\}$$

observed fragments (reads)

$$= \prod_{j=1}^N \sum_{i=1}^M \Pr\{t_i \mid \boldsymbol{\eta}\} \cdot \Pr\{f_j \mid t_i, \mathbf{z}_{ji} = 1\}$$

Prob. of selecting t_i given $\boldsymbol{\eta}$ Prob. of generating fragment f_j given that it originates from t_i

Depends on abundance estimate Independent of abundance estimate

We want to find the values of $\boldsymbol{\eta}$ that **maximize** this probability.
We can do this (at least locally) using an EM algorithm.

A probabilistic view of RNA-Seq quantification

nucleotide fractions known transcriptome assumes independence of fragments

$$\Pr\{\mathcal{F} \mid \boldsymbol{\eta}, \mathcal{T}\} = \prod_{j=1}^N \Pr\{f_j \mid \boldsymbol{\eta}, \mathcal{T}\}$$

observed fragments (reads)

$$= \prod_{j=1}^N \sum_{i=1}^M \Pr\{t_i \mid \boldsymbol{\eta}\} \cdot \Pr\{f_j \mid t_i, \mathbf{z}_{ji} = 1\}$$

Prob. of selecting t_i given $\boldsymbol{\eta}$ Prob. of generating fragment f_j given that it originates from t_i

Depends on abundance estimate Independent of abundance estimate

We want to find the values of $\boldsymbol{\eta}$ that **maximize** this probability.
We can do this (at least locally) using an EM algorithm.

A probabilistic view of RNA-Seq quantification

nucleotide fractions known transcriptome assumes independence of fragments

$$\Pr\{\mathcal{F} \mid \boldsymbol{\eta}, \mathcal{T}\} = \prod_{j=1}^N \Pr\{f_j \mid \boldsymbol{\eta}, \mathcal{T}\}$$

observed fragments (reads)

We can safely truncate $\Pr\{t_i \mid \boldsymbol{\eta}\}$ to 0 for transcripts where a fragment doesn't map/align.

$$= \prod_{j=1}^N \sum_{i=1}^M \Pr\{t_i \mid \boldsymbol{\eta}\} \cdot \Pr\{f_j \mid t_i, z_{ji} = 1\}$$

Prob. of selecting t_i given $\boldsymbol{\eta}$ Prob. of generating fragment f_j given that it originates from t_i

Depends on abundance estimate Independent of abundance estimate

We want to find the values of $\boldsymbol{\eta}$ that **maximize** this probability.
We can do this (at least locally) using an EM algorithm.

A probabilistic view of RNA-Seq quantification

E-step: (what is the “soft assignment” of each read to the transcripts where it aligns)

The diagram illustrates the E-step of RNA-Seq quantification, showing the calculation of the expected value of latent variables given sequenced fragments and parameter estimates at iteration t .

The equation is:

$$E_{Z|f, \eta^{(t)}} = P(Z_{nij} = 1 \mid f, \eta^{(t)}) = \frac{(\eta_i^{(t)} / \ell_i) P(f_n \mid Z_{nij} = 1)}{\sum_{i', j'} (\eta_{i'}^{(t)} / \ell_{i'}) P(f_n \mid Z_{ni'j'} = 1)}$$

Annotations and their corresponding parts of the equation:

- Expected value of latent variables (read $\rightarrow$ txp, pos assignment) given sequenced fragments and parameter estimates at iteration t , $\eta^{(t)}$** : Points to the left-hand side of the equation, $E_{Z|f, \eta^{(t)}}$.
- Current estimate of abundance of txp i** : Points to $\eta_i^{(t)}$ in the numerator.
- Length of txp i** : Points to ℓ_i in the numerator.
- Probability of generating fragment n if it arises from txp i at position j** : Points to $P(f_n \mid Z_{nij} = 1)$ in the numerator.
- Same quantities summed over all reported alignments of fragment n** : Points to the denominator, $\sum_{i', j'} (\eta_{i'}^{(t)} / \ell_{i'}) P(f_n \mid Z_{ni'j'} = 1)$.
- Probability that fragment n truly arises from transcript i at position j** : Points to the entire fraction on the right-hand side of the equation.

A probabilistic view of RNA-Seq quantification

E-step: (what is the “soft assignment” of each read to the transcripts where it aligns)

$$E_{Z|f,\eta^{(t)}} = P(Z_{nij} = 1 \mid f, \eta^{(t)}) = \frac{(\eta_i^{(t)} / \ell_i) P(f_n \mid Z_{nij} = 1)}{\sum_{i',j'} (\eta_{i'}^{(t)} / \ell_{i'}) P(f_n \mid Z_{ni'j'} = 1)}$$

M-step: Given these soft assignments, how abundant is each transcript?

$$\eta_i^{(t+1)} = \frac{E_{Z|f,\eta^{(t)}} [C_i]}{N}$$

Expected number of reads arising from txp i as computed in the E-step


Total number of aligned reads

Estimate of abundance of txp i at iteration t+1

$$\text{where } C_i = \sum_{n,i,j} P(Z_{nij} = 1 \mid f, \eta^{(t)})$$

This approach is quite effective. Unfortunately, it's also quite slow.

Gene expression estimation accuracy in simulated data



What if we model fragment *counts* instead of individual fragments themselves?






Figure 2 from Turro, Ernest, et al. "Haplotype and isoform specific expression estimation using multi-mapping RNA-seq reads." *Genome biology* 12.2 (2011): R13.

Fragment Equivalence Classes



Reads **1** & **3** both map to transcripts B & E


Reads **2** & **4** both map to transcript C

We have 4 reads, but only 2 eq. classes of reads

eq. Label	Count	Aux weights
{B,E}	2	$w^{\{B,E\}}_B, w^{\{B,E\}}_E$
{C}	2	$w^{\{C\}}_C$

This idea goes quite far back in the RNA-seq literature; at least to MMSeq (Turro et al. 2011)

Fragment Equivalence Classes



Reads **1** & **3** both map to transcripts B & E
 Reads **2** & **4** both map to transcript C

w_{ij} encodes the “affinity” of class j to transcript i according to the model. This is $P\{f_j | t_i\}$, aggregated for all fragments in a class.

We have 4 reads, but only 2 eq. classes of reads

eq. Label	Count	Aux weights
{B,E}	2	$w^{\{B,E\}}_B, w^{\{B,E\}}_E$
{C}	2	$w^{\{C\}}_C$

This idea goes quite far back in the RNA-seq literature; at least to MMSeq (Turro et al. 2011)

The number of equivalence classes is **small**

	Yeast	Human	Chicken
# contigs	7353	107,389	335,377
# samples	6	6	8
Total (paired-end) reads	~36,000,000	~116,000,000	~181,402,780
Avg # eq. classes (across samples)	5197	100,535	222,216

The **# of equivalence classes** grows with the complexity of the **transcriptome** — independent of the # of sequence fragments.

Typically, **two or more orders of magnitude** fewer equivalence classes than sequenced fragments.

The offline **inference** algorithm **scales in # of fragment equivalence classes**.

This lets us approximate the likelihood efficiently

Approximate this:

$$\mathcal{L}(\eta; \mathcal{F}) = \prod_{f_j \in \mathcal{F}} \sum_{i=1}^M \Pr(t_i | \eta) \Pr(f_j | t_i)$$

sum over all alignments of fragment

product over all fragments

with this:


$$\mathcal{L}(\eta; \mathcal{F}) \approx \prod_{\mathcal{F}^q \in \mathcal{C}} \left(\sum_{\langle i, t_i \rangle \in \Omega(\mathcal{F}^q)} \Pr(t_i | \eta) \cdot \Pr(f | \mathcal{F}^q, t_i) \right)^{N^q}$$

sum over all transcripts labeling this eq. class


product over all equivalence classes

Why might $\Pr(f_j | t_i)$ matter?

Consider the following scenario:



Conditional probabilities can provide valuable information about origin of a fragment! ***Potentially different for each transcript/fragment pair.***



Prob of observing a fragment of size ~200 is **large**


Prob of observing a fragment of size ~450 is **small**

**Many terms can be considered in a general “fragment-transcript agreement” model¹.
e.g. position, orientation, alignment path etc.**

¹ “Salmon provides fast and bias-aware quantification of transcript expression”, Nature Methods 2017

Why might $\Pr(f_j | t_i)$ matter?

Consider the following scenario:



We'll return to this question later

Conditional probabilities can provide valuable information about origin of a fragment! *Potentially different for each transcript/fragment pair.*


Prob of observing a fragment of size ~200 is **large**

Prob of observing a fragment of size ~450 is **small**

Many terms can be considered in a general “fragment-transcript agreement” model¹.
e.g. position, orientation, alignment path etc.

¹ “Salmon provides fast and bias-aware quantification of transcript expression”, Nature Methods 2017

Actual RNA-seq protocols are a bit more “involved”



There is **substantial** potential for biases and deviations from the *basic* model — indeed, we see quite a few.

Biases abound in RNA-seq data

Biases in prep & sequencing
can have a significant effect on the
fragments we see:

Fragment gc-bias¹—


The GC-content of the fragment
affects the likelihood of sequencing

Sequence-specific bias²—

sequences surrounding fragment
affect the likelihood of sequencing

Positional bias²—

fragments sequenced non-uniformly
across the body of a transcript



1:Love, Michael I., John B. Hogenesch, and Rafael A. Irizarry. "Modeling of RNA-seq fragment sequence bias reduces systematic errors in transcript abundance estimation." bioRxiv (2015): 025767.

2:Roberts, Adam, et al. "Improving RNA-Seq expression estimates by correcting for fragment bias." Genome biology 12.3 (2011): 1.

Biases abound in RNA-seq data

Basic idea (1): Modify the “effective length” of a transcript to account for changes in the sampling probability. This leads to changes in soft-assignment in EM -> changes in TPM.

Basic idea (2): The effective length of a transcript is the sum of the bias terms at each position across a transcript. The bias term at a given position is simply the (observed / expected) sampling probability.

The trick is how to define “expected” given only biased data.

Bias Modeling

Bias correction works by adjusting the effective lengths of the transcripts:
The effective length becomes the sum of the per-base biases

$$\tilde{\ell}'_i = \sum_{j=1}^{j \leq \ell_i} \sum_{k=1}^{k \leq f_i(j,L)} \frac{b_{gc+}(t_i, j, j+k)}{b_{gc-}(t_i, j, j+k)} \cdot \frac{b_{s+}^{5'}(t_i, j)}{b_{s-}^{5'}(t_i, j)} \cdot \frac{b_{s+}^{3'}(t_i, j+k)}{b_{s-}^{3'}(t_i, j+k)} \cdot \frac{b_{p+}^{5'}(t_i, j+k)}{b_{p-}^{5'}(t_i, j+k)} \cdot \frac{b_{p+}^{3'}(t_i, j+k)}{b_{p-}^{3'}(t_i, j+k)} \cdot \Pr\{X = j\}$$

Fragment GC bias model:


Density of fragments with specific GC content,
conditioned on GC fraction at read start/end

Foreground:


Observed

Background:

Expected given est. abundances



GC-fraction of fragment



Bias Modeling

Bias correction works by adjusting the effective lengths of the transcripts:
The effective length becomes the sum of the per-base biases

$$\tilde{\ell}'_i = \sum_{j=1}^{j \leq \ell_i} \sum_{k=1}^{k \leq f_i(j,L)} \frac{b_{gc+}(t_i, j, j+k)}{b_{gc-}(t_i, j, j+k)} \cdot \frac{b_{s+}^{5'}(t_i, j)}{b_{s-}^{5'}(t_i, j)} \cdot \frac{b_{s+}^{3'}(t_i, j+k)}{b_{s-}^{3'}(t_i, j+k)} \cdot \frac{b_{p+}^{5'}(t_i, j+k)}{b_{p-}^{5'}(t_i, j+k)} \cdot \frac{b_{p+}^{3'}(t_i, j+k)}{b_{p-}^{3'}(t_i, j+k)} \cdot \Pr\{X = j\}$$

Seq-specific bias model*:

VLMM for the 10bp window surrounding the 5' read start site and the 3' read start site

Foreground:

Observed

Background:


Expected given est. abundances



Add this sequence to training set with weight =
 $P\{f \mid t_i\}$

Same, but independent
model for 3' end

Priming bias is sample & sequence-specific



Basic idea

The sequencing is unbiased w.r.t. sequence context if

$$\begin{array}{ccc} & E[x_i | s_i] = N \Pr[m_i | s_i] = N \Pr[m_i] = E[x_i] & \\ \nearrow & & \nwarrow \\ \text{Expected read count at pos } i & = & \text{Unconditional expected read count at pos } i \\ \text{conditioned on sequence} & & \end{array}$$

Define the sequence bias as: $b_i = \Pr[s_i] / \Pr[s_i | m_i]$

So that:

$$E[b_i x_i | s_i] = b_i E[x_i | s_i] = N b_i \Pr[m_i | s_i] = N \frac{\Pr[m_i | s_i] \Pr[s_i]}{\Pr[s_i | m_i]} = N \Pr[m_i] = E[x_i]$$

Priming bias is sample & sequence-specific





Table 3. The Pearson’s correlation coefficient r between log-adjusted read counts and log-adjusted TaqMan values

Method	Correlation
Unadjusted	0.6650**
7mer	0.6680**
GLM	0.6874**
MART	0.6998*
BN	0.7086

The best *model* may also be sample-specific



Contrast with Roberts et al. which uses a fixed-structure
VLMM to model the sample-specific bias.

Bias Modeling

Bias correction works by adjusting the effective lengths of the transcripts:
The effective length becomes the sum of the per-base biases


$$\tilde{\ell}'_i = \sum_{j=1}^{j \leq \ell_i} \sum_{k=1}^{k \leq f_i(j,L)} \frac{b_{gc+}(t_i, j, j+k)}{b_{gc-}(t_i, j, j+k)} \cdot \frac{b_{s+}^{5'}(t_i, j)}{b_{s-}^{5'}(t_i, j)} \cdot \frac{b_{s+}^{3'}(t_i, j+k)}{b_{s-}^{3'}(t_i, j+k)} \cdot \frac{b_{p+}^{5'}(t_i, j+k)}{b_{p-}^{5'}(t_i, j+k)} \cdot \frac{b_{p+}^{3'}(t_i, j+k)}{b_{p-}^{3'}(t_i, j+k)} \cdot \Pr\{X = j\}$$

Position bias model*:

Foreground:
Observed


Density of 5' and 3' read start positions —
different models for transcripts of different length

Background:
Expected given est. abundances




*Roberts, Adam, et al. "Improving RNA-Seq expression estimates by correcting for fragment bias." Genome biology 12.3 (2011): 1.

Salmon provides fast and bias-aware quantification of transcript expression




Patro, R., Duggal, G., Love, M. I., Irizarry, R. A., & Kingsford, C. (2017).
Salmon provides fast and bias-aware quantification of transcript expression. *Nature Methods*.

Salmon's "pipeline"



Phase 1: Online Inference (asynchronous!)



Compute local η' using η^{t-1} & current "bias" model to allocate fragments

Update global nucleotide fractions: $\eta^t = \eta^{t-1} + a^t \eta'$

Update "bias" model

Weighting factor that
decays over time

Place mappings in **equivalence classes**

- Have access to **all fragment-level information** when making these updates
- Often converges very quickly.
- Compare-And-Swap (CAS) for synchronizing updates of different batches

*Based on: Foulds et al. Stochastic collapsed variational Bayesian inference for latent Dirichlet allocation. ACM SIGKDD, 2013.

Broderick, Tamara, et al. "Streaming variational bayes." *Advances in Neural Information Processing Systems*. 2013.

Hsieh, Cho-Jui, Hsiang-Fu Yu, and Inderjit S. Dhillon. "PASSCoDe: Parallel ASynchronous Stochastic dual Co-ordinate Descent." *ICML*. Vol. 15. 2015.

Raman, Parameswaran, et al. "Extreme Stochastic Variational Inference: Distributed and Asynchronous." *arXiv preprint arXiv:1605.09499* (2016). (@ICML 2017)

Give each transcript appropriate prior mass η^0 (init.)

For each mini-batch B^t of reads {

For each read r in B^t {

For each alignment a of r {

compute (un-normalized) prob of a using η^{t-1} , and aux params

}

normalize alignment probs & update local transcript weights η'


add / update the equivalence class for read r

sample $a \in r$ to update auxiliary models

}

update global transcript weights given local transcript weights according to "update rule" $\Rightarrow \eta^t = \eta^{t-1} + w^t \eta'$

}



mini-batches processed in parallel by different threads

additive nature of updates mitigates effects of
no synchronization between mini-batches

Broderick, Tamara, et al. "Streaming variational bayes." *Advances in Neural Information Processing Systems*. 2013.


Hsieh, Cho-Jui, Hsiang-Fu Yu, and Inderjit S. Dhillon. "PASSCoDe: Parallel ASynchronous Stochastic dual Co-ordinate Descent." *ICML*. Vol. 15. 2015.

Raman, Parameswaran, et al. "Extreme Stochastic Variational Inference: Distributed and Asynchronous." *arXiv preprint arXiv:1605.09499* (2016). (@ICML 2017)

In this phase, we maintain *current* estimates of abundance.

Each group of fragments arrive (streaming), and we use their mapping locations & current estimates to:

1. Allocate them to transcripts
2. Update auxiliary models
3. Place them in **equivalence classes**




online inference
[SCVB0]

We use a streaming, parallel, stochastic inference algorithm for Phase 1; a variant of **Stochastic Collapsed Variational Bayesian Inference** [SCVB0]*

* Foulds, James, et al. "Stochastic collapsed variational Bayesian inference for latent Dirichlet allocation." Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining. ACM, 2013.

Optimizing the objective



our ML objective has a simple, **closed-form update rule** in terms of our eq. classes

$$\alpha_i^{u+1} = \sum_{\mathcal{F}^q \in \mathcal{C}} N^q \left(\frac{\alpha_i^u w_i^q}{\sum_{\langle k, t_k \rangle \in \Omega(\mathcal{F}^q)} \alpha_k^u w_k^q} \right)$$

count of eq. class j

weight of t_i in eq. class q

estimated read count from transcript i at iteration u+1

$$\hat{\eta}_i = \frac{\alpha_i}{\sum_j \alpha_j}$$


we also provide the *option* to use a **variational Bayesian** objective instead

Accuracy difference can be larger with biased data

Simulated data:
2 conditions; 8 samples each

- Simulated transcripts across entire genome with known abundance using Polyester (modified to account for GC bias)
- How well do we recover the underlying relative abundances?
- How does accuracy vary with level of bias?

Lower is better



Sequence-bias models don't account for fragment-level GC bias


Mis-estimates confound downstream analysis

Simulated data:

2 conditions; 8 replicates each

- set 10% of txps to have fold change of 1/2 or 2 — rest unchanged.
- How well do we recover true DE?
- Since bias is systematic, effect may be even worse than accuracy difference suggests.

Recovery of DE transcripts



Importance with **experimental** data

30 samples from the GEUVADIS study:

15 samples from UNIGE sequencing center

15 samples from CNAG_CRG sequencing center

Same human population, expect few-to-no *real* DE

Randomized condition assignments result in $\ll 1$ DE txp

DE of data between centers (FDR < 1%) (TPM > 0.1)

	Salmon	Kallisto	eXpress
All transcripts	1,183	2,620	2,472
Transcripts of 2 isoform genes	228	545	531

Bias and batch effects are *substantial*, and must be accounted for.

Importance with **experimental** data

30 samples from the GEUVADIS study:

15 samples from UNIGE sequencing center

15 samples from CNAG_CRG sequencing center

Effects seem **at least as extreme** at the gene level

DE of data between centers (FDR < 1%) (TPM > 0.1)

	Salmon	Kallisto	eXpress
All genes	455	1,200	1,582
Transcripts of 2 isoform genes	224	545	531

Bias and **batch effects** are *substantial*, and must be accounted for.

Further improving the factorization (at low computational cost)

Bioinformatics, 33, 2017, i142–i151
doi: 10.1093/bioinformatics/btx262
ISMB/ECCB 2017



Improved data-driven likelihood factorizations for transcript abundance estimation

Mohsen Zakeri, Avi Srivastava, Fatemeh Almodaresi and Rob Patro*


Department of Computer Science, Stony Brook University, Stony Brook, NY 11790, USA


A probabilistic view of RNA-Seq quantification

We want to find the values of η that **maximize** this probability.
We can do this (at least locally) using the EM algorithm.

but

This leads to an iterative EM algorithm where *each iteration* scales in the total number of **alignments** in the sample (typically on the order of 10^7 — 10^8), and typically 10^2 — 10^3 **iterations**


$$\mathcal{L}(\eta; \mathcal{F}, \mathcal{T}) = \prod_{f \in \mathcal{F}} \sum_{t_i \in \Omega(f)} \Pr(t_i \mid \eta) \Pr(f \mid t_i)$$




Set of transcripts where f maps/aligns

Recall : Fragment Equivalence Classes

$$f \sim f' \iff \Omega(f) = \Omega(f')$$

$$\Omega(f) = \{t \mid f \text{ maps to } t\}$$



Reads 1 & 3 both map to transcripts B & E

Reads 2 & 4 both map to transcript C


We have 4 reads, but only 2 eq. classes/types of reads

eq. Label	Count
{B,E}	2
{C}	2

Equivalence classes in RNA-Seq quantification

Long history of this idea — collapsing “redundant” reads

This list is not-complete (just illustrative)



This lets us approximate the likelihood efficiently

Approximate this:

$$\mathcal{L}(\boldsymbol{\eta}; \mathcal{F}, \mathcal{T}) = \prod_{f \in \mathcal{F}} \sum_{t_i \in \Omega(f)} \Pr(t_i \mid \boldsymbol{\eta}) \Pr(f \mid t_i)$$

sum over all alignments of fragment

product over all fragments

with this:

$$\mathcal{L}(\boldsymbol{\eta}; \mathcal{F}, \mathcal{T}) \approx \prod_{\mathcal{F}^q \in \mathcal{C}} \left(\sum_{t_i \in \Omega(\mathcal{F}^q)} \Pr(t_i \mid \boldsymbol{\eta}) \cdot \Pr(f \mid \mathcal{F}^q, t_i) \right)^{N^q}$$

sum over all transcripts labeling this eq. class


product over all equivalence classes

The approximation applies because **all** f in $\mathcal{F}^q$ have **the same**

conditional probability given t_i — i.e. $\Pr(f \mid \mathcal{F}^q, t_i)$

Why might $\Pr(f_j | t_i)$ matter?

Consider the following scenario:




Conditional probabilities can provide valuable information about origin of a fragment! **Potentially different for each transcript/fragment pair.**

Prob of observing a fragment of size ~ 200 is **large**
Prob of observing a fragment of size ~ 450 is **small**

**Many terms can be considered in a general “fragment-transcript agreement” model¹.
e.g. position, orientation, alignment path etc.**


¹ “Salmon provides fast and bias-aware quantification of transcript expression”, Nature Methods 2017

Does this term matter?



- Transcripts of RAD51 gene — txp coverage drawn randomly in [1,200]
- Distribution over 30 random replicates of this distribution

Does this term matter?



- Transcripts of RAD51 gene — txp coverage drawn randomly in [1,200]
- Distribution over 30 random replicates of this distribution

Range-factorized equivalence relation

Recall:

$$f \sim f' \iff \Omega(f) = \Omega(f')$$

$$\Omega(f) = \{t \mid f \text{ maps to } t\}$$

Now:

$b_k(f, \langle t_{i_1}, \dots, t_{i_j} \rangle)$ Given a fragment and vector of transcripts, returns a vector of bin indices — each in $[0, k)$ — that encode the conditional bin into which f falls with respect to each transcript.

of conditional bins. Default = $4 + \lceil \sqrt{|\Omega(\mathcal{F}^q)|} \rceil$


$$f \sim_r f' \iff \Omega(f) = \Omega(f') \wedge b_k(f, \Omega(f)) = b_k(f', \Omega(f'))$$

Maps to the same set of transcripts

Has the same *binned* cond. prob vector


Range-based factorization

60 fragments in equivalence class $\{t_1, t_2\}$




Range-based factorization improves approximation

60 fragments in equivalence class $\{t_1, t_2\}$



- Provides a way to control the divergence between the full and factorized conditional likelihood distributions of an equivalence class

How well does this work?



- Transcripts of RAD51 gene — txp coverage drawn randomly in [1,200]
- Distribution over 30 random replicates of this distribution

Transcriptome-wide assessment can mask important differences

- Over tens of thousands of transcripts — overall differences are small
- But, we know this; factorized approaches are known to work well generally^{1,2,3,4}

	Method	MARD	Spearman
~ factorization	<i>Salmon-U</i>	0.24	0.80
	<i>Salmon</i>	0.22	0.81
~ _r factorization	<i>Salmon-RF</i>	0.21	0.83
no factorization	<i>Salmon-FM</i>	0.21	0.83
	<i>eXpress</i>	0.29	0.78
	<i>eXpress (+50)</i>	0.23	0.83
	<i>RSEM</i>	0.21	0.82

- 30M paired-end reads, simulated with RSEM-Sim

1) Turro, Ernest, et al. "Haplotype and isoform specific expression estimation using multi-mapping RNA-seq reads." *Genome biology* 12.2 (2011): R13.

2) Srivastava, Avi, et al. "RapMap: a rapid, sensitive and accurate tool for mapping RNA-seq reads to transcriptomes." *Bioinformatics* 32.12 (2016): i192-i200.

3) Bray, N. L., et al. "Near-optimal probabilistic RNA-seq quantification." *Nature biotechnology* 34.5 (2016): 525.

4) Patro, Rob, et al. "Salmon provides fast and bias-aware quantification of transcript expression." *Nature Methods* 14.4 (2017): 417-419.

Transcriptome-wide assessment can mask important differences

- Focus on a subset of “critical” transcripts (not too easy, not intractable)
- Transcripts where RSEM yields an ARD in [0.25,0.75]

	Method	MARD	Spearman
~ factorization	<i>Salmon-U</i>	0.46	0.56
	<i>Salmon</i>	0.43	0.58
~ _r factorization	<i>Salmon-RF</i>	0.41	0.64
no factorization	<i>Salmon-FM</i>	0.41	0.65
	<i>eXpress</i>	0.53	0.54
	<i>eXpress (+50)</i>	0.48	0.59
	<i>RSEM</i>	0.41	0.65

- 30M paired-end reads, simulated with RSEM-Sim


1) Turro, Ernest, et al. "Haplotype and isoform specific expression estimation using multi-mapping RNA-seq reads." *Genome biology* 12.2 (2011): R13.

2) Srivastava, Avi, et al. "RapMap: a rapid, sensitive and accurate tool for mapping RNA-seq reads to transcriptomes." *Bioinformatics* 32.12 (2016): i192-i200.

3) Bray, N. L., et al. "Near-optimal probabilistic RNA-seq quantification." *Nature biotechnology* 34.5 (2016): 525.

4) Patro, Rob, et al. "Salmon provides fast and bias-aware quantification of transcript expression." *Nature Methods* 14.4 (2017): 417-419.

Range-factorization improves correlation with full-model on experimental data



SEQC samples from UHRR (SRR1215996 - SRR1217002)

7 technical replicates to define distribution

Treat RSEM results as **ground truth** (though clearly, it's not perfect)

Range-factorization is still very (computationally) efficient

Factorization “size” on simulated data


	<i>Salmon-U</i>	<i>Salmon</i>	<i>Salmon-RF</i>	<i>Salmon-FM</i>
# eq. classes	438,393	438,393	625,638	29,447,710
# hits	5,986,371	5,986,371	8,212,669	103,663,423

eq. classes : The number of different “types” of read — i.e. $\sum_{\mathcal{F}^q \in \mathcal{C}} 1$


hits : The number of hits is the sum, over each equivalence class, of the number of transcripts in this equivalence class — i.e. $\sum_{\mathcal{F}^q \in \mathcal{C}} |\Omega(\mathcal{F}^q)|$

Difference is *marginal* with respect to # of reads / alignments

Range-factorization is still very (computationally) efficient



Range-factorization controls memory requirements



Estimating Posterior Uncertainty

One “issue” with maximum likelihood (ML)

The generative statistical model is a principled and elegant way to represent the RNA-seq process.

It can be optimized efficiently using e.g. the EM / VBEM algorithm.


but, these efficient optimization algorithms return “point estimates” of the abundances. That is, there is no notion of how *certain* we are in the computed abundance of transcript.

One “issue” with maximum likelihood (ML)


There are multiple sources of uncertainty e.g.

- Technical variance : If we sequenced the *exact* same sample again, we’d get a different set of fragments, and, potentially a different solution.
- Uncertainty in inference: We are almost never guaranteed to find a unique, globally optimal result. If we started our algorithm with different initialization parameters, we might get a different result.

We’re trying to find the *best* parameters in a space with 10s to 100s of thousands of dimensions!



One “issue” with maximum likelihood (ML)



Assessing Uncertainty

There are a few ways to address this “issue”

Do a fully Bayesian inference¹:

Infer the entire posterior distribution of parameters, not just a ML estimate (e.g. using MCMC) — too slow!

✓ Posterior Gibbs Sampling^{2,3}:

Starting from our ML estimate, do MCMC sampling to explore how parameters vary — if our ML estimate is good, this can be made *quite fast*.

✓ Bootstrap Sampling⁴:

Resample (from range-factorized equivalence class counts) with replacement, and re-run the ML estimate for each sample. This can be made reasonably fast.

Happy to discuss details / implications of this further.

1: BitSeq (with MCMC) actually does this. It's very accurate, but very slow. [Glaus, Peter, Antti Honkela, and Magnus Rattray. "Identifying differentially expressed transcripts from RNA-seq data with biological variation." *Bioinformatics* 28.13 (2012): 1721-1728.]

2: RSEM has the ability to do this, and it seems to work well, but each sample scales in the # of reads. [Li, Bo, and Colin N. Dewey. "RSEM: accurate transcript quantification from RNA-Seq data with or without a reference genome." *BMC bioinformatics* 12.1 (2011): 1.]

3: MMSEQ can perform Gibbs sampling over shared variables (i.e. equiv classes), producing estimates from the mean of the posterior dist. Turro, Ernest, et al. "Haplotype and isoform specific expression estimation using multi-mapping RNA-seq reads." *Genome biology* 12.2 (2011): 1.

4: IsoDE introduced the idea of bootstrapping counts to assess quantification uncertainty. [Al Seesi, Sahar, et al. "Bootstrap-based differential gene expression analysis for RNA-Seq data with and without replicates." *BMC genomics* 15.8 (2014): 1.], but it was first made practical / fast in kallisto by doing the bootstrapping over equivalence classes.

A few ways to implement Gibbs Sampling for this problem

The model of MMSeq

$$X_{it} \mid \mu_t \sim \text{Pois}(bs_i M_{it} \mu_t), \quad (12)$$

$$\mu_t \sim \text{Gam}(\alpha, \beta). \quad (13)$$

The full conditionals are:

$$\{X_{i1}, \dots, X_{it}\} \mid \{\mu_1, \dots, \mu_t\}, k_i \sim \text{Mult} \left(k_i, \frac{M_{i1}\mu_1}{\sum_t M_{it}\mu_t}, \dots, \frac{M_{in}\mu_n}{\sum_t M_{it}\mu_t} \right), \quad (14)$$

$$\mu_t \mid \{X_{1t}, \dots, X_{mt}\} \sim \text{Gam} \left(\alpha + \sum_i X_{it}, \beta + bl_t \right). \quad (15)$$

Again, the s_i are not needed as they are absent from the full conditionals.

This uncertainty matters




Figure 2.10: **Posterior distribution of expression levels of three transcripts of gene Q6ZMZ0.** The posterior distribution is represented in form of a histogram of expression samples converted into Log RPKM expression measure. The dashed lines mark the mean expression for each transcript.

This uncertainty matters






Figure 2.12: **Exon model of transcripts of gene Q6ZMZ0.** (a) transcript sequence profile obtained from the UCSC genome browser (Kuhn et al., 2013). In this annotation, transcript uc001bwm.3 has different 3' untranslated region and transcript uc010oho.1 has extra nucleotides at the end of second exon. As the second change cannot be distinguished in the UCSC genome browser diagram, we provide schematic splice variant model highlighting the differences (b).



*Glaus, Peter. *Bayesian Methods for Gene Expression Analysis from High-throughput Sequencing Data*. Diss. University of Manchester, 2014.

This uncertainty matters

We observe considerably increased variance due to read mapping ambiguity



If we know this increased uncertainty, we can propagate it & use it in downstream analysis (differential expression)!